

Peter Racioppo

Los Angeles, CA · (424) 273-0644 · pcracioppo@gmail.com · linkedin.com/in/peter-racioppo · github.com/PCR-git

Machine Learning Engineer and researcher specializing in Transformer architectures, sequence modeling, and dynamical systems. Solo author of an ICML 2026 Spotlight paper and first author on multiple machine learning publications, with experience training language models, developing novel attention mechanisms, and building training/eval infrastructure.

EXPERIENCE

Research Assistant | Center for Vision, Cognition, Learning, and Autonomy (UCLA) 2023 – 2025

- Developed novel Transformer architectures grounded in dynamical systems, Bayesian filtering, and state estimation, culminating in a solo-authored ICML 2026 Spotlight paper.
- Built PyTorch training infrastructure and conducted large-scale language-model experiments to evaluate long-context extrapolation and architectural design choices.

Machine Learning Research Intern | Doheny Eye Institute 2024 – 2026

- Developed a knowledge distillation framework for retinal image segmentation across multiple ocular diseases.
- Developed a novel Sliding Window Attention mechanism that imposes temporal stationarity priors for robust learning in low-data regimes; validated on longitudinal retinal disease forecasting.
- First author on two peer-reviewed medical AI publications. ([MDPI Bioengineering 2025](#), [MDPI Life 2026](#)).

RESEARCH HIGHLIGHTS

Robust Filter Attention: Self-Attention as Precision-Weighted State Estimation. P. Racioppo. [[Paper](#) | [GitHub](#)]

International Conference on Machine Learning (ICML, 2026) – Spotlight

- Solo-authored ICML 2026 Spotlight paper, deriving a new class of attention mechanisms from state-estimation principles and demonstrating improved language-model perplexity relative to RoPE across long-context benchmarks.
- Built and trained Transformer language models end-to-end, conducting extensive ablations and maintaining stable zero-shot extrapolation at up to 8× training context length.
- Open-sourced implementation with reproducible experiments.

The Transformer as a Polar State Estimator. P. Racioppo. [[Preprint](#)]

- Developed a probabilistic geometric framework interpreting Transformer attention as approximate state estimation on the hypersphere, unifying attention, residual connections, and normalization.
- Proposed the *Polar Transformer*, a geometry-aware Transformer architecture incorporating intrinsic manifold updates and magnitude-dependent attention precision.

SELECTED PROJECTS

Hallucination Detection for LLMs. [[Preprint](#) | [GitHub](#)]

- Developed a correctness-estimation system that predicts LLM answer reliability from internal Transformer activations.
- Built an adaptive retrieval framework that selectively triggers and evaluates RAG based on predicted correctness.

AC-SINDy: Compositional Sparse Identification of Nonlinear Dynamics. [[Preprint](#) | [GitHub](#)]

- Developed a scalable method for learning symbolic models of dynamical systems with improved robustness to noise.
- Implemented models and training pipeline and evaluated on simulated dynamical systems.

SKILLS

Languages: Python, R, MATLAB, SQL.

Core Areas: Sequence modeling, NLP, computer vision, dynamical systems, deep learning, statistical modeling.

ML & Deep Learning: PyTorch, scikit-learn, Transformers, CNNs, PEFT/QLoRA, model training/finetuning, evaluation.

Systems & Tools: Git, GPU training, Hugging Face, cloud training, Pandas, NumPy, Jupyter.

EDUCATION

MS in Statistics, UCLA — 2025

MS in Mechanical Engineering (Controls), UCLA — 2023

BS in Physics, UCLA